

Universal Sentence Encoder on Amazon Reviews

By Alex Vaillant

Research Question

"To what extent can we accurately identify our future customer's sentiment in unseen reviews based on the past reviews on our products by using NLP and NN techniques? The end model will be used as part of an early "warning" system to trigger the need for intervention and praise based on the review's sentiment."

Set Up Environment

```
In [1]: from platform import python_version
        print(python_version())
```

3.7.10

Load Libraries

```
In [2]: import pandas as pd
        import numpy as np
        import tensorflow as tf
        import keras
        from keras import layers, Sequential
        from tensorflow.keras.layers import Dense
        from tensorflow.keras.models import load_model
        import tensorflow_hub as hub
        import tensorflow_text
        from sklearn.preprocessing import OneHotEncoder
        from sklearn.model_selection import train_test_split
```

Load Universal Sentence Encoder

```
In [3]: USE = hub.load("https://tfhub.dev/google/universal-sentence-encoder-multilingual-large/
```

Data Extraction

```
In [4]: URL = 'C:/Users/tedda/Desktop/Data Science Portfolio/Machine Learning/Sentiment Analysis/
amazon_data = pd.read_csv(URL, delimiter = '\t', names = ['Review', 'Review_Type'])
print('Amazon Raw Data Shape:', amazon_data.shape)
```

Amazon Raw Data Shape: (1000, 2)

Data Preparation

One Hot Encode the Review Labels

```
In [5]: onehot = OneHotEncoder(sparse=False).fit_transform(amazon_data['Review_Type'].to_numpy(
```

```
In [6]: onehot
```

```
Out[6]: array([[1., 0.],
               [0., 1.],
               [0., 1.],
               ...,
               [1., 0.],
               [1., 0.],
               [1., 0.]])
```

Split data into Training and Testing sets (80%, 20%)

```
In [7]: x_train, x_test, y_train, y_test = train_test_split(amazon_data['Review'], onehot, test
```

USE the Reviews' Text (X)

```
In [8]: X_train = []
        for review in x_train:
            embedded = USE(review)
            reviews_embedded = tf.reshape(embedded, [-1]).numpy()
            X_train.append(reviews_embedded)

        X_train = np.array(X_train)
```

```
In [9]: X_train.shape
```

```
Out[9]: (800, 512)
```

```
In [10]: X_test = []
          for review in x_test:
              embedded = USE(review)
              reviews_embedded = tf.reshape(embedded, [-1]).numpy()
              X_test.append(reviews_embedded)

          X_test = np.array(X_test)
```

```
In [11]: X_test.shape
```

```
Out[11]: (200, 512)
```

Export Cleansed Datasets

```
In [12]: pd.DataFrame(X_train).to_csv('C:/Users/tedda/Desktop/Data Science Portfolio/Machine Lea
pd.DataFrame(X_test).to_csv('C:/Users/tedda/Desktop/Data Science Portfolio/Machine Lear
pd.DataFrame(y_train).to_csv('C:/Users/tedda/Desktop/Data Science Portfolio/Machine Lea
pd.DataFrame(y_test).to_csv('C:/Users/tedda/Desktop/Data Science Portfolio/Machine Lear
```

Build the Model

```
In [13]: model = keras.Sequential()

model.add(keras.layers.Dense(256, input_shape=(X_train.shape[1], ), activation = 'relu')

model.add(keras.layers.Dropout(rate=0.5))

model.add(keras.layers.Dense(128, activation = 'relu'))

model.add(keras.layers.Dropout(rate=0.5))

model.add(keras.layers.Dense(2, activation = 'softmax'))

model.compile(loss='categorical_crossentropy', optimizer = 'adam', metrics = ['accuracy'])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	131328
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 2)	258
Total params: 164,482		
Trainable params: 164,482		
Non-trainable params: 0		

Train the Model

```
In [14]: model.fit(X_train, y_train, epochs = 4, batch_size = 20, validation_split = 0.2, verbose
```

```
Epoch 1/4
32/32 [=====] - 40s 36ms/step - loss: 0.6711 - accuracy: 0.6349
- val_loss: 0.5480 - val_accuracy: 0.8125
Epoch 2/4
32/32 [=====] - ETA: 0s - loss: 0.4438 - accuracy: 0.88 - 0s 6ms
s/step - loss: 0.4337 - accuracy: 0.8861 - val_loss: 0.3684 - val_accuracy: 0.8625
Epoch 3/4
32/32 [=====] - 0s 6ms/step - loss: 0.2444 - accuracy: 0.9297 -
val_loss: 0.3626 - val_accuracy: 0.8687
Epoch 4/4
32/32 [=====] - 0s 6ms/step - loss: 0.1853 - accuracy: 0.9414 -
val_loss: 0.3911 - val_accuracy: 0.8438
```

```
Out[14]: <keras.callbacks.History at 0x265f3f1b4c8>
```

Save and Load the Model

```
In [15]: model_url = 'C:/Users/tedda/Desktop/Data Science Portfolio/Machine Learning/Sentiment A
model.save(model_url)
```

```
SA_model = load_model(model_url)
```

Evaluate the Model's Accuracy

In [16]:

```
SA_model.evaluate(X_test, y_test)
```

```
7/7 [=====] - 1s 4ms/step - loss: 0.1862 - accuracy: 0.9250
```

Out[16]: [0.18624836206436157, 0.925000011920929]